

Big Data Infrastructure

CS 489/698 Big Data Infrastructure (Winter 2016)

Week 7: Analyzing Relational Data (2/3)

February 23, 2016

Jimmy Lin

David R. Cheriton School of Computer Science

University of Waterloo

These slides are available at <http://lintool.github.io/bigdata-2016w/>

This work is licensed under a Creative Commons Attribution-Noncommercial-Share Alike 3.0 United States
See <http://creativecommons.org/licenses/by-nc-sa/3.0/us/> for details




facebook®

Jeff Hammerbacher, Information Platforms and the Rise of the Data Scientist.
In, *Beautiful Data*, O'Reilly, 2009.

“On the first day of logging the Facebook clickstream, more than 400 gigabytes of data was collected. The load, index, and aggregation processes for this data set really taxed the Oracle data warehouse. Even after significant tuning, we were unable to aggregate a day of clickstream data in less than 24 hours.”

SQL-on-Hadoop



What not just use a
database to begin with?
Cost + Scalability

Databases are great...

If your data has structure (and you know what the structure is)

If your data is reasonably clean

If you know what queries you're going to run ahead of time

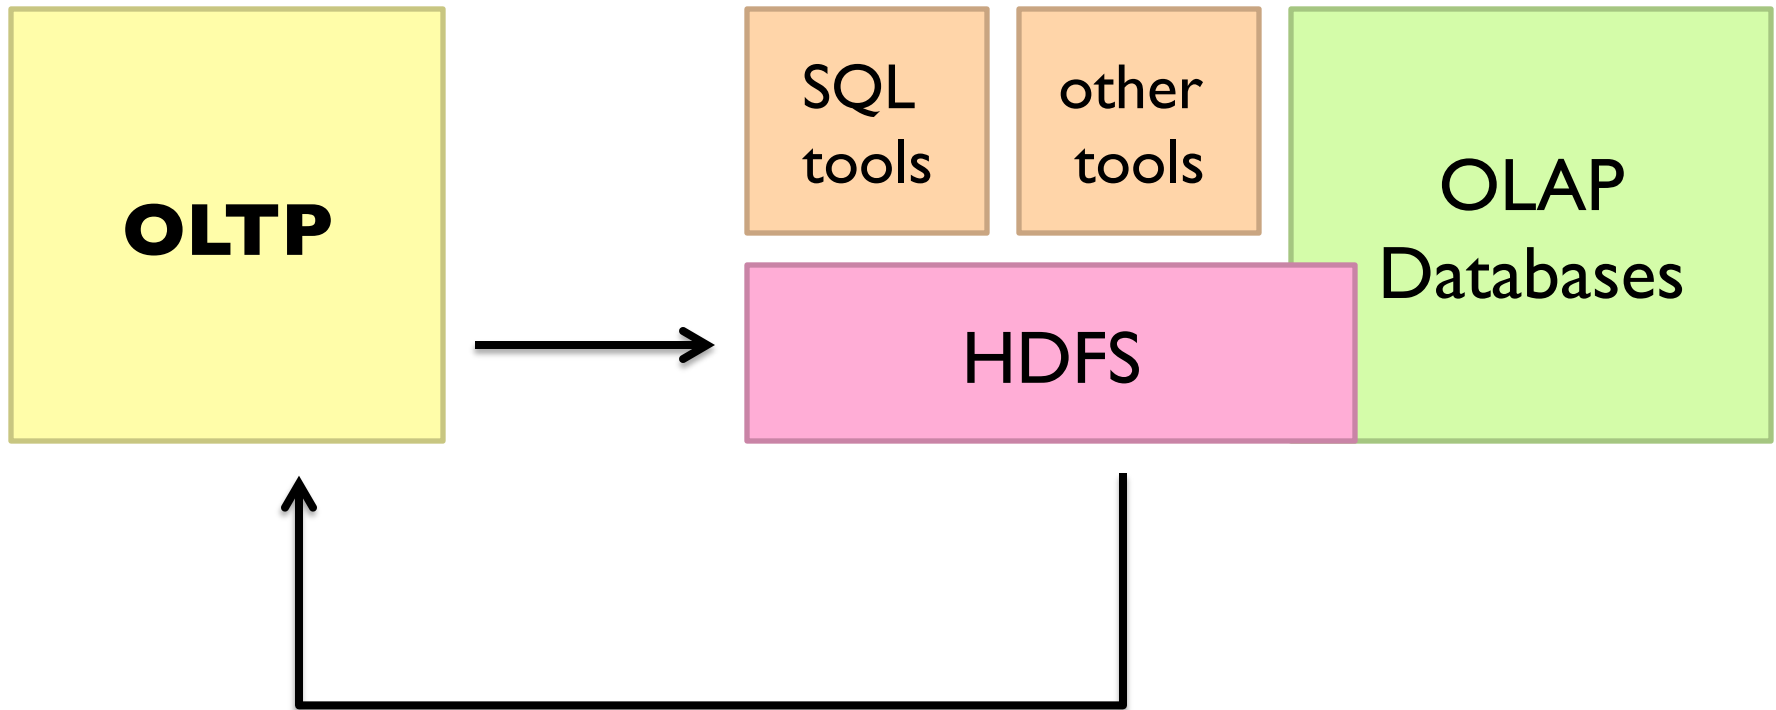
Databases are not so great...

If your data has little structure (or you don't know the structure)


If your data is messy and noisy

If you don't know what you're looking for

What's the selling point of SQL-on-Hadoop?
Trade (a little?) performance for flexibility



SQL-on-Hadoop



Hive: Example


- Relational join on two tables:
 - Table of word counts from Shakespeare collection
 - Table of word counts from the bible

```
SELECT s.word, s.freq, k.freq FROM shakespear s
JOIN bible k ON (s.word = k.word) WHERE s.freq >= 1 AND k.freq >= 1
ORDER BY s.freq DESC LIMIT 10;
```

the	25848	62394
I	23031	8854
and	19671	38985
to	18038	13526
of	16700	34654
a	14170	8057
you	12702	2720
my	11297	4135
in	10797	12445
is	8882	6884


Hive: Behind the Scenes

```
SELECT s.word, s.freq, k.freq FROM shakespeare s  
JOIN bible k ON (s.word = k.word) WHERE s.freq >= 1 AND k.freq >= 1  
ORDER BY s.freq DESC LIMIT 10;
```



(Abstract Syntax Tree)

```
(TOK_QUERY (TOK_FROM (TOK_JOIN (TOK_TABREF shakespeare s) (TOK_TABREF bible k) (= (. (TOK_TABLE_OR_COL s)  
word) (. (TOK_TABLE_OR_COL k) word)))) (TOK_INSERT (TOK_DESTINATION (TOK_DIR TOK_TMP_FILE)) (TOK_SELECT  
(TOK_SELEXPR (. (TOK_TABLE_OR_COL s) word)) (TOK_SELEXPR (. (TOK_TABLE_OR_COL s) freq)) (TOK_SELEXPR (.  
(TOK_TABLE_OR_COL k) freq))) (TOK_WHERE (AND (>= (. (TOK_TABLE_OR_COL s) freq) 1) (>= (. (TOK_TABLE_OR_COL k)  
freq) 1))) (TOK_ORDERBY (TOK_TABSORTCOLNAMEDESC (. (TOK_TABLE_OR_COL s) freq))) (TOK_LIMIT 10)))
```



(one or more of MapReduce jobs)

Hive: Behind the Scenes

STAGE DEPENDENCIES:

Stage-1 is a root stage
Stage-2 depends on stages: Stage-1
Stage-0 is a root stage

STAGE PLANS:

Stage: Stage-1
Map Reduce

Alias -> Map Operator Tree:

```
s
  TableScan
  alias: s
  Filter Operator
  predicate:
    expr: (freq >= 1)
    type: boolean
  Reduce Output Operator
  key expressions:
    expr: word
    type: string
  sort order: +
  Map-reduce partition columns:
    expr: word
    type: string
  tag: 0
  value expressions:
    expr: freq
    type: int
    expr: word
    type: string
```

```
k
  TableScan
  alias: k
  Filter Operator
  predicate:
    expr: (freq >= 1)
    type: boolean
  Reduce Output Operator
  key expressions:
    expr: word
    type: string
  sort order: +
  Map-reduce partition columns:
    expr: word
    type: string
  tag: 1
  value expressions:
    expr: freq
    type: int
```

Reduce Operator Tree:

```
Join Operator
condition map:
  Inner Join 0 to 1
condition expressions:
  0 {VALUE._col0} {VALUE._col1}
  1 {VALUE._col0}
outputColumnNames: _col0, _col1, _col2
Filter Operator
predicate:
  expr: ((_col0 >= 1) and (_col2 >= 1))
  type: boolean
Select Operator
expressions:
  expr: _col1
  type: string
  expr: _col0
  type: int
  expr: _col2
  type: int
outputColumnNames: _col0, _col1, _col2
File Output Operator
compressed: false
GlobalTableId: 0
table:
  input format: org.apache.hadoop.mapred.SequenceFileInputFormat
  output format: org.apache.hadoop.hive ql.io.HiveSequenceFileOutputFormat
```

Stage: Stage-2

Map Reduce

Alias -> Map Operator Tree:

hdfs://localhost:8022/tmp/hive-training/364214370/10002

Reduce Output Operator

key expressions:

expr: _col1

type: int

sort order: -

tag: -1

value expressions:

expr: _col0

type: string

expr: _col1

type: int

expr: _col2

type: int

Reduce Operator Tree:

Extract

Limit

File Output Operator

compressed: false

GlobalTableId: 0

table:

input format: org.apache.hadoop.mapred.TextInputFormat

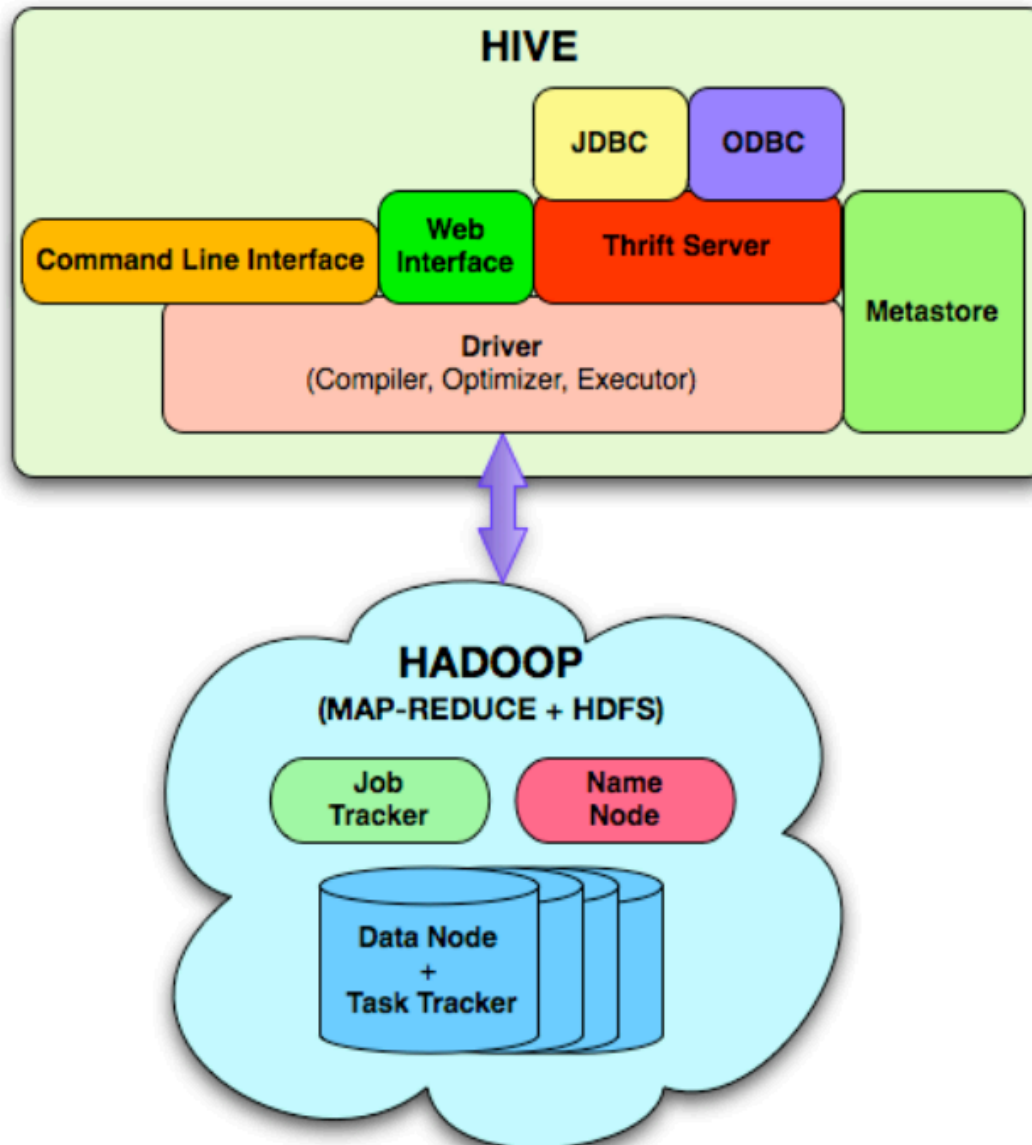
output format: org.apache.hadoop.hive ql.io.HiveIgnoreKeyTextOutputFormat

Stage: Stage-0

Fetch Operator

limit: 10

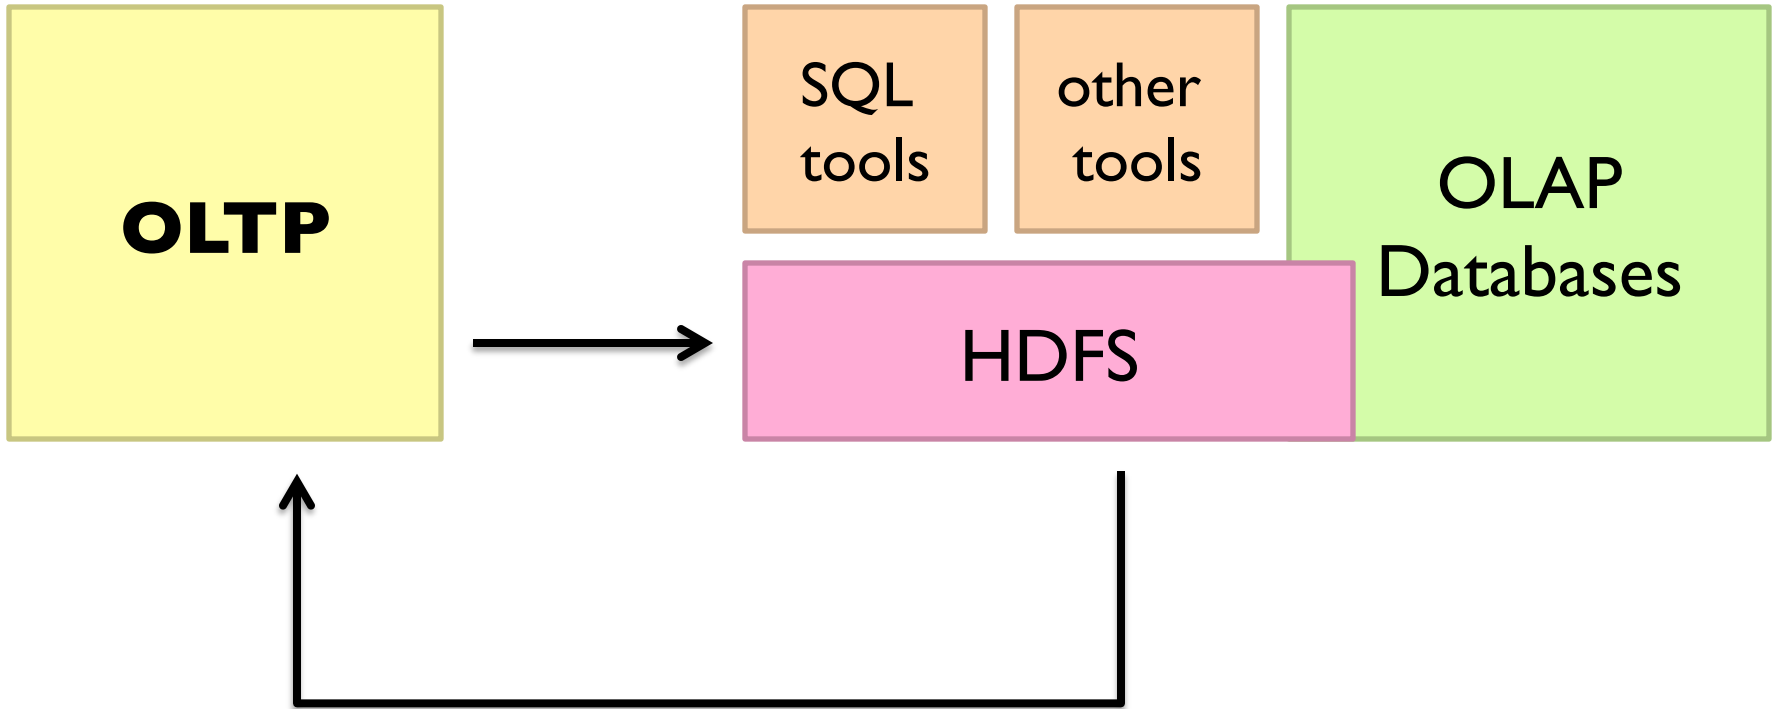
Hive Architecture



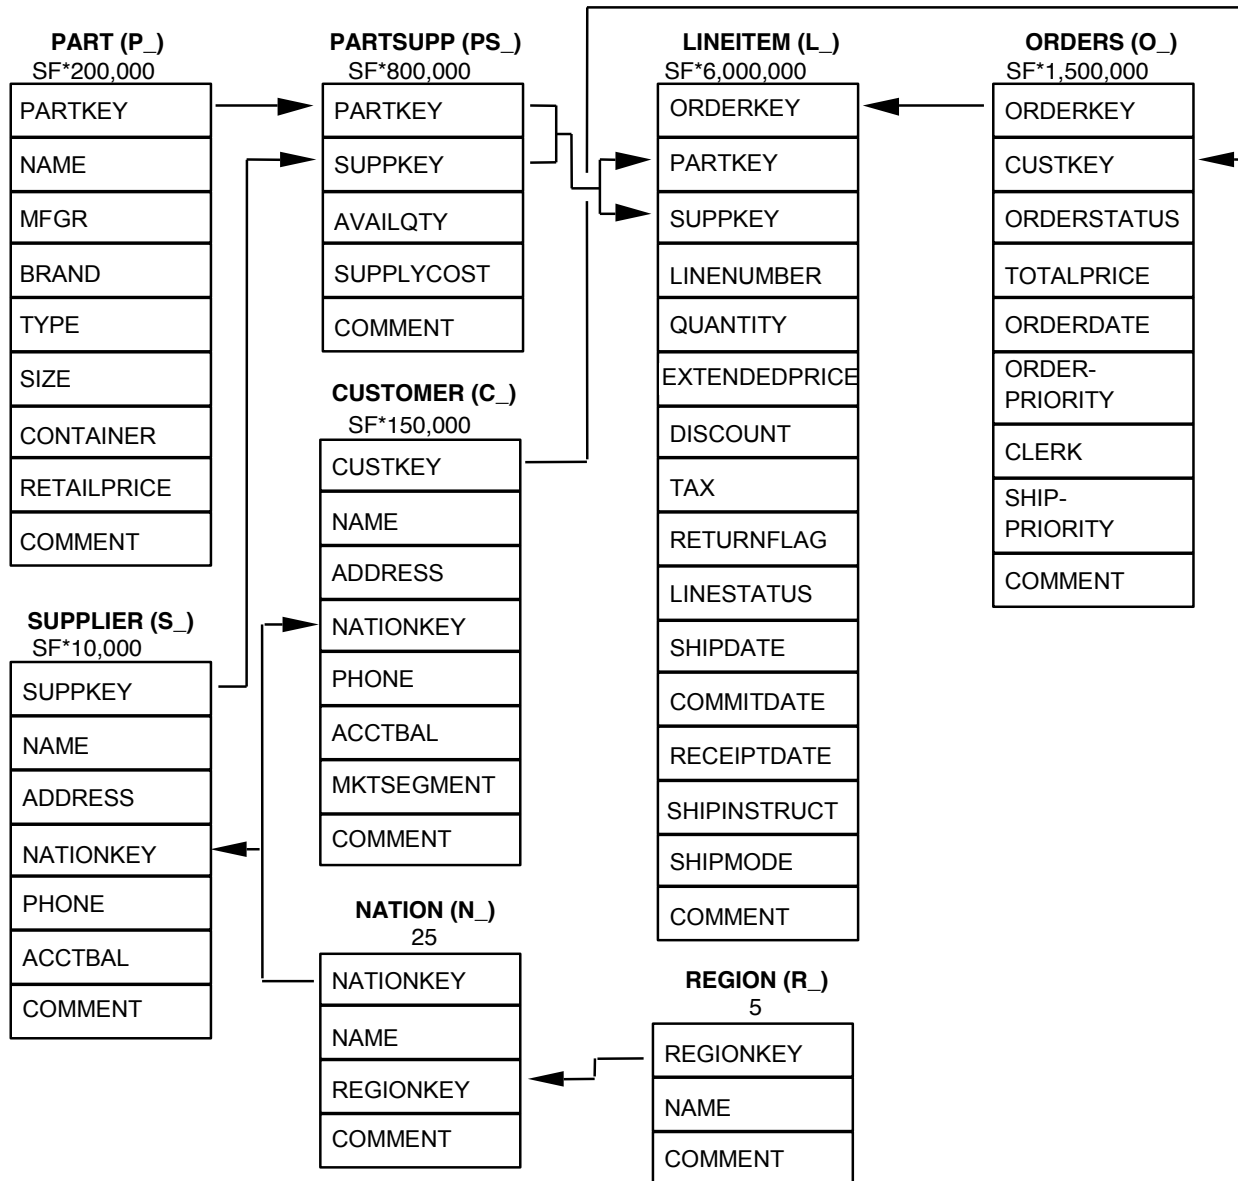
Hive Implementation

- Metastore holds metadata
 - Databases, tables
 - Schemas (field names, field types, etc.)
 - Permission information (roles and users)
- Hive data stored in HDFS
 - Tables in directories
 - Partitions of tables in sub-directories
 - Actual data in files (plain text or binary encoded)

Feature or bug?
(this is the essence of SQL-on-Hadoop)



TPC-H Data Warehouse



A black and white photograph of a garden bed. A border made of bricks is visible, with some bricks having square holes. Various plants, including a small tree and some leafy shrubs, are growing in the garden bed. The text "MapReduce algorithms for processing relational data" is overlaid in the center of the image.

MapReduce algorithms for processing relational data

Relational Algebra


○ Primitives

- Projection (π)
- Selection (σ)
- Cartesian product (\times)
- Set union (\cup)
- Set difference ($-$)
- Rename (ρ)

○ Other operations

- Join (\bowtie)
- Group by... aggregation
- ...


Selection




Selection in MapReduce

- Easy!
 - In mapper: process each tuple, only emit tuples that meet criteria
 - Can be pipelined with projection
 - No reducers necessary (unless to do something else)
- Performance mostly limited by HDFS throughput
 - Speed of encoding/decoding tuples becomes important
 - Take advantage of compression when available
 - Semistructured data? No problem!

Projection



$\xrightarrow{\pi_{\square\bigcirc}}$



Projection in MapReduce

- Easy!

- In mapper: process each tuple, re-emit with only projected attributes
- Can be pipelined with selection
- No reducers necessary (unless to do something else)

- Implementation detail: bookkeeping required

- Need to keep track of attribute mappings after projection
e.g., name was `r[4]`, becomes `r[1]` after projection

- Performance mostly limited by HDFS throughput

- Speed of encoding/decoding tuples becomes important
- Take advantage of compression when available
- Semistructured data? No problem!

Group by... Aggregation

- Aggregation functions:
 - AVG
 - MAX
 - MIN
 - SUM
 - COUNT
 - ...
- MapReduce implementation:
 - Map over dataset, emit tuples, keyed by group by attribute
 - Framework automatically groups values by group by attribute
 - Compute aggregation function in reducer
 - Optimize with combiners, in-mapper combining

You already know how to do this!

Remember this!
(week 2)

Combiner Design

- Combiners and reducers share same method signature
 - Sometimes, reducers can serve as combiners
 - Often, not...
- Remember: combiner are optional optimizations
 - Should not affect algorithm correctness
 - May be run 0, 1, or multiple times
- Example: find average of integers associated with the same key

Remember this?
(week 2)

SELECT key, AVG(value) FROM r GROUP BY key;

Combiner Design

- Combiner
 - Some
 - Often
- Remember
 - Should
 - May be
- Example

Computing the Mean: Version 1

Computing the Mean: Version 2

```
1: class  
2: me  
3:  
4:  
5:  
6:  
7:  
8:  
9:  
1: class  
2: me  
3:  
4:  
5:  
6:  
7:  
8:  
9:
```

Computing the Mean: Version 3

```
1: class  
2: me  
3:  
4:  
5:  
6:  
7:  
8:  
9:  
1: class  
2: me  
3:  
4:  
5:  
6:  
7:  
8:  
9:
```

Computing the Mean: Version 4

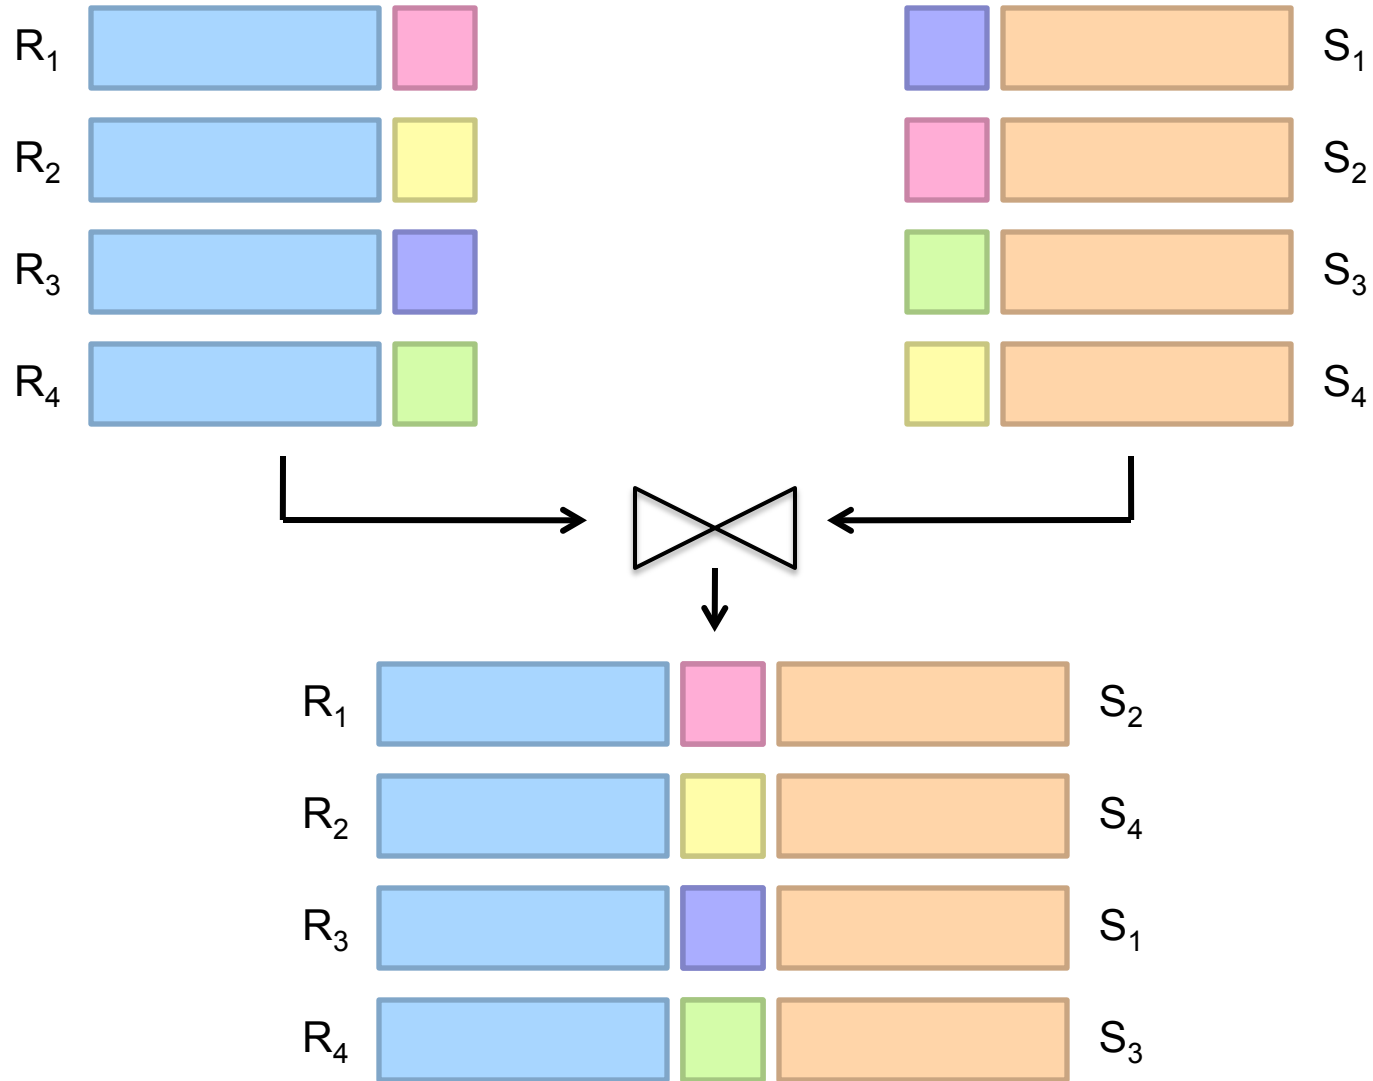
```
1: class MAPPER  
2:   method INITIALIZE  
3:      $S \leftarrow \text{new ASSOCIATIVEARRAY}$   
4:      $C \leftarrow \text{new ASSOCIATIVEARRAY}$   
5:   method MAP(string  $t$ , integer  $r$ )  
6:      $S\{t\} \leftarrow S\{t\} + r$   
7:      $C\{t\} \leftarrow C\{t\} + 1$   
8:   method CLOSE  
9:     for all term  $t \in S$  do  
10:       EMIT(term  $t$ , pair ( $S\{t\}, C\{t\}$ ))
```

Are combiners still needed?


Relational Joins




Relational Joins




Types of Relationships



Many-to-Many



One-to-Many



One-to-One

Join Algorithms in MapReduce


- Reduce-side join
 - aka repartition join
 - aka shuffle join
- Map-side join
 - aka sort-merge join
- Hash join
 - aka broadcast join
 - aka replicated join

Reduce-side Join aka repartition join, shuffle join

- Basic idea: group by join key
 - Map over both datasets
 - Emit tuple as value with join key as the intermediate key
 - Execution framework brings together tuples sharing the same key
 - Perform join in reducer
- Two variants
 - 1-to-1 joins
 - 1-to-many and many-to-many joins


Reduce-side Join: 1-to-1

Map



Remember to “tag” the tuple as being from R or S ...


Reduce



Note: no guarantee if R is going to come first or S

Reduce-side Join: 1-to-many

Map



Reduce



What's the problem?

Quick Aside: Secondary Sorting

- MapReduce sorts input to reducers by key
 - Values are arbitrarily ordered
- What if want to sort value also?
 - E.g., $k \rightarrow (v_1, r_1), (v_3, r_2), (v_4, r_3), (v_8, r_4) \dots$

Secondary Sorting: Solutions

○ Solution 1:

- Buffer values in memory, then sort
- Why is this a bad idea?

○ Solution 2:

- “Value-to-key conversion” design pattern:
form composite intermediate key, (k, v_i)
- Let execution framework do the sorting
- Preserve state across multiple key-value pairs to handle processing
- Anything else we need to do?

Value-to-Key Conversion

Before

$k \rightarrow (v_8, r_4), (v_1, r_1), (v_4, r_3), (v_3, r_2) \dots$

Values arrive in arbitrary order...

After

$(k, v_1) \rightarrow r_1$

$(k, v_3) \rightarrow r_2$

$(k, v_4) \rightarrow r_3$

$(k, v_8) \rightarrow r_4$

...

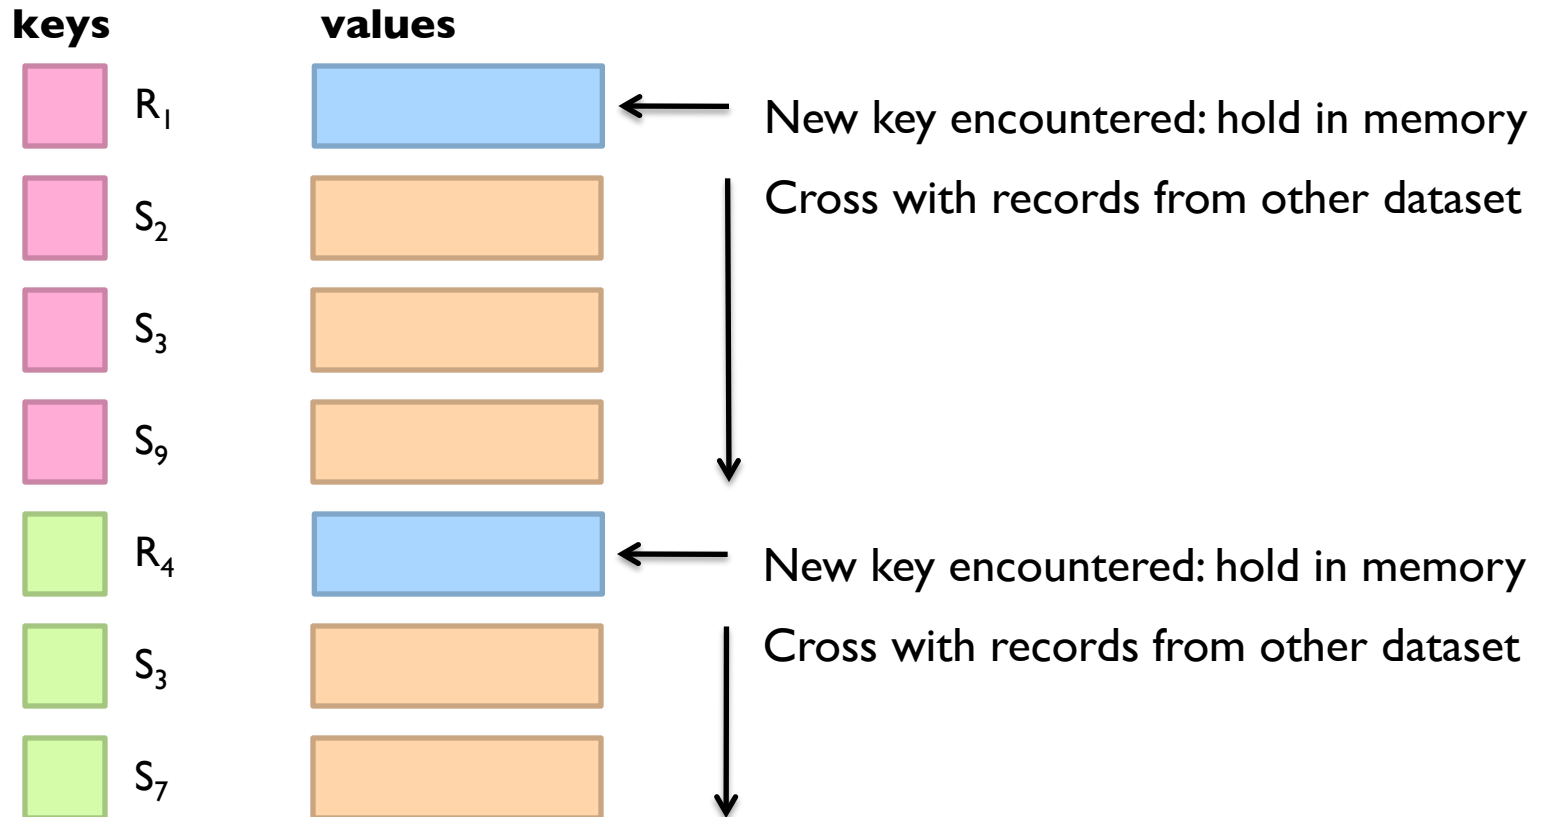
Values arrive in sorted order...

Process by preserving state across multiple keys

Remember to partition correctly!

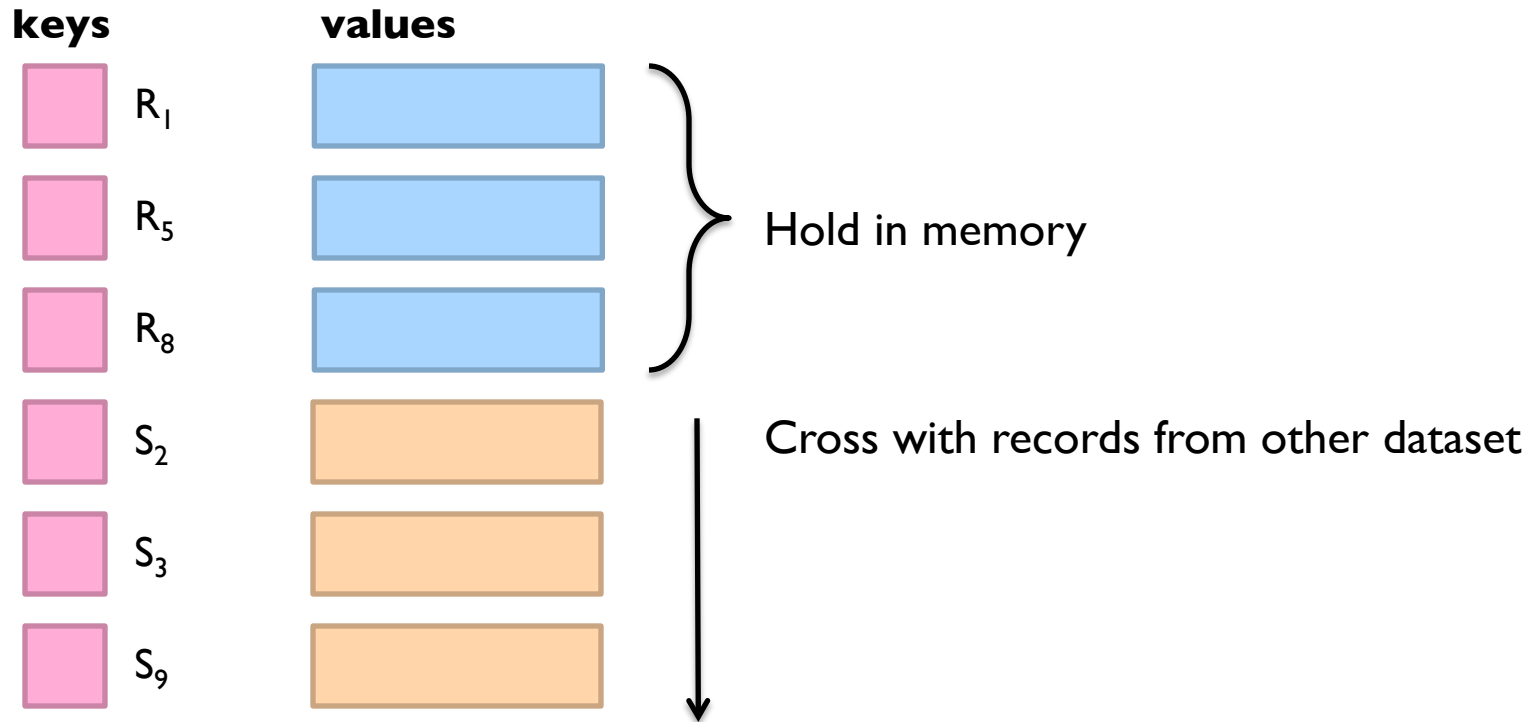
Reduce-side Join: V-to-K Conversion

In reducer...



Reduce-side Join: many-to-many


In reducer...



What's the problem?


Map-side Join aka sort-merge join

Assume two datasets are sorted by the join key:



Map-side Join aka sort-merge join

Assume two datasets are sorted by the join key:



How can we parallelize this? Co-partitioning

Map-side Join aka sort-merge join

- Works if...
 - Two datasets are co-partitioned
 - Sorted by join key
- MapReduce implementation:
 - Map over one dataset, read from other corresponding partition
 - No reducers necessary (unless to do something else)
- Co-partitioned, sorted datasets: realistic to expect?

Hash Join aka broadcast join, replicated join

- Basic idea:
 - Load one dataset into memory in a hashmap, keyed by join key
 - Read other dataset, probe for join key
- Works if...
 - $R \ll S$ and R fits into memory
- MapReduce implementation:
 - Distribute R to all nodes (e.g., DistributedCache)
 - Map over S, each mapper loads R in memory and builds the hashmap
 - For every tuple in S, probe join key in R
 - No reducers necessary (unless to do something else)


Hash Join Variants

- Co-partitioned variant:
 - R and S co-partitioned (but not sorted)?
 - Only need to build hashmap on the corresponding partition
- Striped variant:
 - R too big to fit into memory?
 - Divide R into R_1, R_2, R_3, \dots s.t. each R_n fits into memory
 - Perform hash join: $\forall n, R_n \bowtie S$
 - Take the union of all join results
- Use a global key-value store:
 - Load R into memcached (or Redis)
 - Probe global key-value store for join key

Which join to use?

- In-memory join > map-side join > reduce-side join
- Limitations of each?
 - In-memory join: memory
 - Map-side join: sort order and partitioning
 - Reduce-side join: general purpose

SQL-on-Hadoop



Putting Everything Together

```
SELECT big1.fx, big2.fy, small.fz
FROM big1
JOIN big2 ON big1.id1 = big2.id1
JOIN small ON big1.id2 = small.id2
WHERE big1.fx = 2015 AND
       big2.f1 < 40 AND
       big2.f2 > 2;
```

Build logical plan

Optimize logical plan

Select physical plan


Putting Everything Together

```
SELECT big1.fx, big2.fy, small.fz
FROM big1
JOIN big2 ON big1.id1 = big2.id1
JOIN small ON big1.id2 = small.id2
WHERE big1.fx = 2015 AND
      big2.f1 < 40 AND
      big2.f2 > 2;
```

Build logical plan

Optimize logical plan

Select physical plan




Putting Everything Together

```
SELECT big1.fx, big2.fy, small.fz
FROM big1
JOIN big2 ON big1.id1 = big2.id1
JOIN small ON big1.id2 = small.id2
WHERE big1.fx = 2015 AND
       big2.f1 < 40 AND
       big2.f2 > 2;
```

Build logical plan

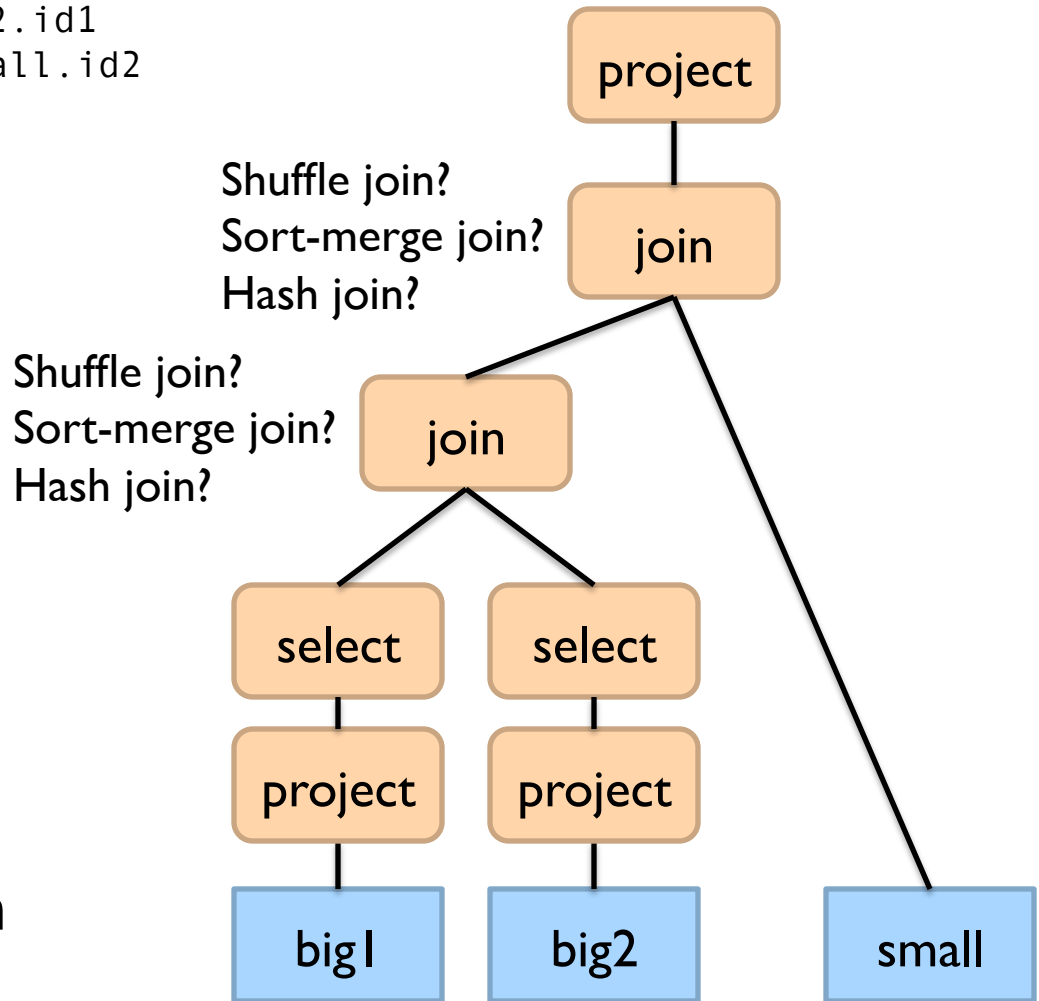
Optimize logical plan

Select physical plan



Putting Everything Together

```
SELECT big1.fx, big2.fy, small.fz
FROM big1
JOIN big2 ON big1.id1 = big2.id1
JOIN small ON big1.id2 = small.id2
WHERE big1.fx = 2015 AND
       big2.f1 < 40 AND
       big2.f2 > 2;
```



Build logical plan

Optimize logical plan

Select physical plan


Putting Everything Together

```
SELECT big1.fx, big2.fy, small.fz
FROM big1
JOIN big2 ON big1.id1 = big2.id1
JOIN small ON big1.id2 = small.id2
WHERE big1.fx = 2015 AND
      big2.f1 < 40 AND
      big2.f2 > 2;
```

Build logical plan

Optimize logical plan

Select physical plan



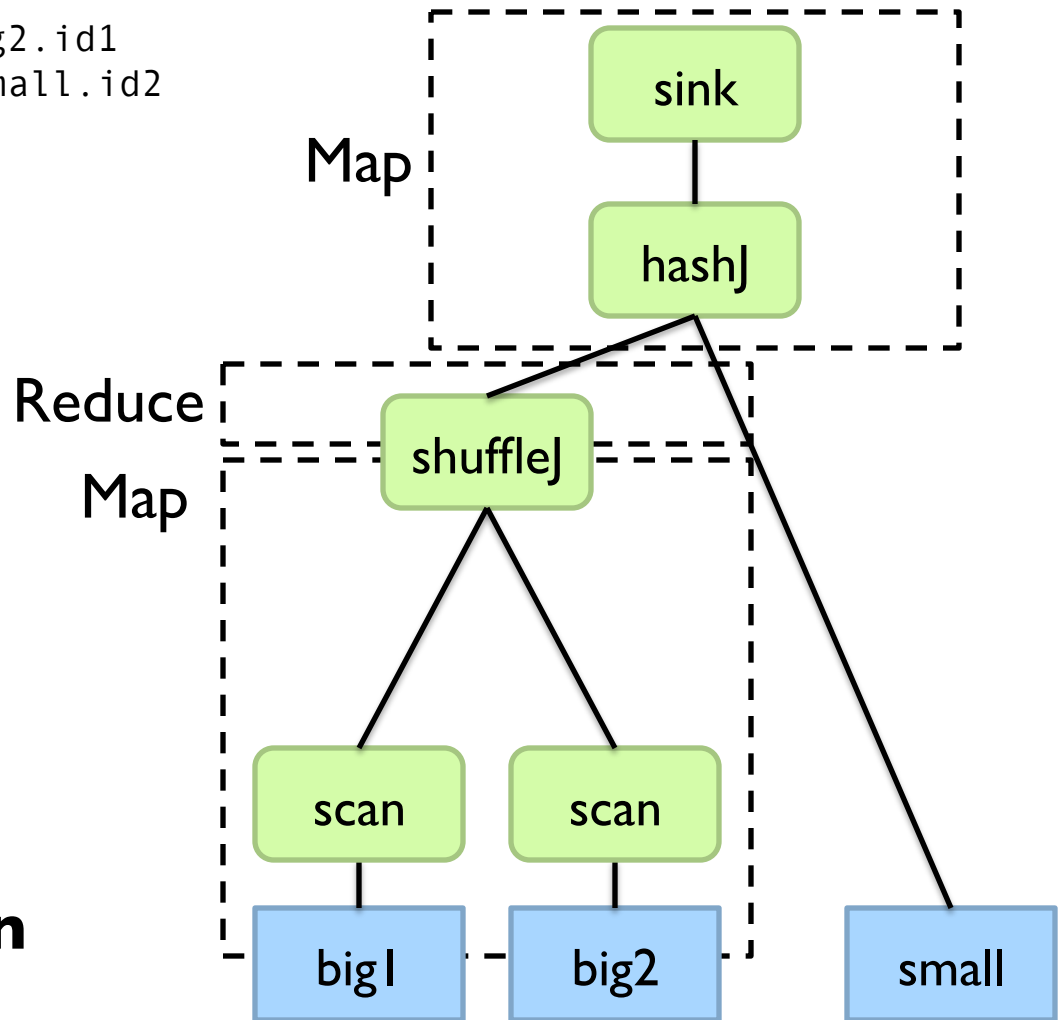
Putting Everything Together

```
SELECT big1.fx, big2.fy, small.fz
FROM big1
JOIN big2 ON big1.id1 = big2.id1
JOIN small ON big1.id2 = small.id2
WHERE big1.fx = 2015 AND
       big2.f1 < 40 AND
       big2.f2 > 2;
```

Build logical plan

Optimize logical plan


Select physical plan



Hive: Behind the Scenes


**Now you understand
what's going on here!**

```
SELECT s.word, s.freq, k.freq FROM shakespear s  
JOIN bible k ON (s.word = k.word) WHERE s.freq >= 1 AND k.freq >= 1  
ORDER BY s.freq DESC LIMIT 10;
```



(Abstract Syntax Tree)

```
(TOK_QUERY (TOK_FROM (TOK_JOIN (TOK_TABREF shakespear s) (TOK_TABREF bible k) (= (. (TOK_TABLE_OR_COL s)  
word) (. (TOK_TABLE_OR_COL k) word)))) (TOK_INSERT (TOK_DESTINATION (TOK_DIR TOK_TMP_FILE)) (TOK_SELECT  
(TOK_SELEXPR (. (TOK_TABLE_OR_COL s) word)) (TOK_SELEXPR (. (TOK_TABLE_OR_COL s) freq)) (TOK_SELEXPR (.  
(TOK_TABLE_OR_COL k) freq))) (TOK_WHERE (AND (>= (. (TOK_TABLE_OR_COL s) freq) 1) (>= (. (TOK_TABLE_OR_COL k)  
freq) 1))) (TOK_ORDERBY (TOK_TABSORTCOLNAMEDESC (. (TOK_TABLE_OR_COL s) freq))) (TOK_LIMIT 10)))
```



(one or more of MapReduce jobs)

Hive: Behind the Scenes

**Now you understand
what's going on here!**

STAGE DEPENDENCIES:

Stage-1 is a root stage
Stage-2 depends on stages: Stage-1
Stage-0 is a root stage

STAGE PLANS:

Stage: Stage-1

Map Reduce

Alias -> Map Operator Tree:

```
s
  TableScan
  alias: s
  Filter Operator
  predicate:
    expr: (freq >= 1)
    type: boolean
  Reduce Output Operator
  key expressions:
    expr: word
    type: string
  sort order: +
  Map-reduce partition columns:
    expr: word
    type: string
  tag: 0
  value expressions:
    expr: freq
    type: int
    expr: word
    type: string
```

```
k
  TableScan
  alias: k
  Filter Operator
  predicate:
    expr: (freq >= 1)
    type: boolean
  Reduce Output Operator
  key expressions:
    expr: word
    type: string
  sort order: +
  Map-reduce partition columns:
    expr: word
    type: string
  tag: 1
  value expressions:
    expr: freq
    type: int
```

Reduce Operator Tree:

```
Join Operator
condition map:
  Inner Join 0 to 1
condition expressions:
  0 {VALUE._col0} {VALUE._col1}
  1 {VALUE._col0}
outputColumnNames: _col0, _col1, _col2
Filter Operator
predicate:
  expr: ((_col0 >= 1) and (_col2 >= 1))
  type: boolean
Select Operator
expressions:
  expr: _col1
  type: string
  expr: _col0
  type: int
  expr: _col2
  type: int
outputColumnNames: _col0, _col1, _col2
File Output Operator
compressed: false
GlobalTableId: 0
table:
  input format: org.apache.hadoop.mapred.SequenceFileInputFormat
  output format: org.apache.hadoop.hive ql.io.HiveSequenceFileOutputFormat
```

Stage: Stage-2

Map Reduce

Alias -> Map Operator Tree:

hdfs://localhost:8022/tmp/hive-training/364214370/10002

Reduce Output Operator

key expressions:

expr: _col1

type: int

sort order: -

tag: -1

value expressions:

expr: _col0

type: string

expr: _col1

type: int

expr: _col2

type: int

Reduce Operator Tree:

Extract

Limit

File Output Operator

compressed: false

GlobalTableId: 0

table:

input format: org.apache.hadoop.mapred.TextInputFormat


output format: org.apache.hadoop.hive ql.io.HiveIgnoreKeyTextOutputFormat

Stage: Stage-0

Fetch Operator

limit: 10

SQL-on-Hadoop



What about Spark SQL?

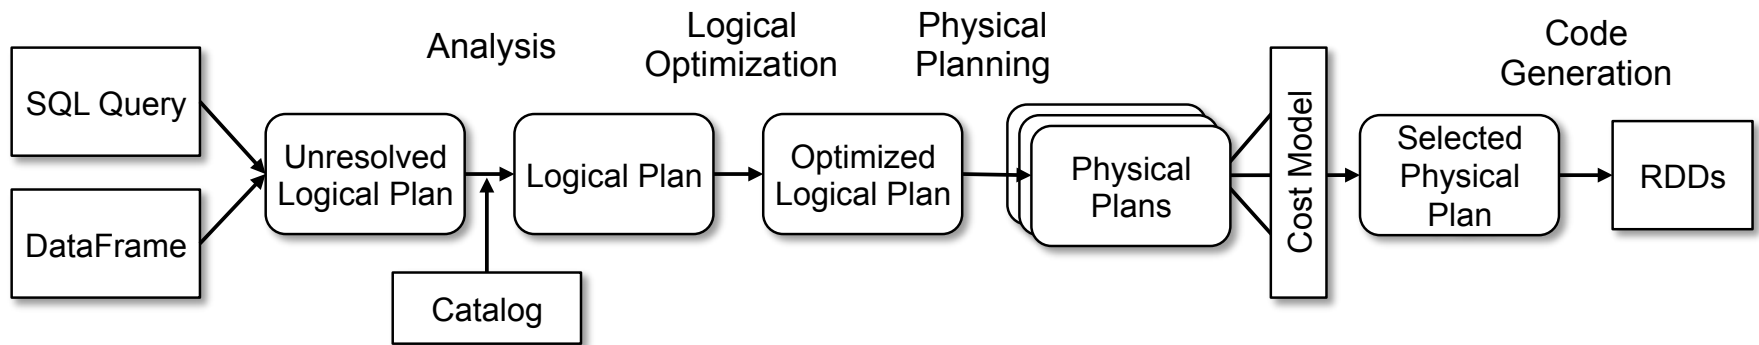
- Based on the DataFrame API:
 - A distributed collection of data organized into named columns
- Two ways of specifying SQL queries:
 - Directly:

```
val sqlContext = ... // An existing SQLContext
val df = sqlContext.sql("SELECT * FROM table")
// df is a dataframe, can be further manipulated...
```

- Via DataFrame API:

```
// employees is a dataframe:
employees
  .join(dept, employees("deptId") === dept("id"))
  .where(employees("gender") === "female")
  .groupBy(dept("id"), dept("name"))
  .agg(count("name"))
```


Spark SQL: Query Planning




At the end of the day... it's transformations on RDDs

Spark SQL: Physical Execution


Narrow Dependencies:



map, filter




union




join with inputs
co-partitioned

= Map-side join

Wide Dependencies:



groupByKey



join with inputs not
co-partitioned


= Reduce-side join

Hash join with broadcast variables

Hadoop Data Warehouse Design

- Observation:
 - Joins are relatively expensive
 - OLAP queries frequently involve joins
- Solution: denormalize
 - What's normalization again?
 - Why normalize to begin with?
 - Fundamentally a time-space tradeoff
 - How much to denormalize?
 - What about consistency?


Denormalization Opportunities?



“Denormalizing the snowflake”


What's the assignment?

SQL-on-Hadoop

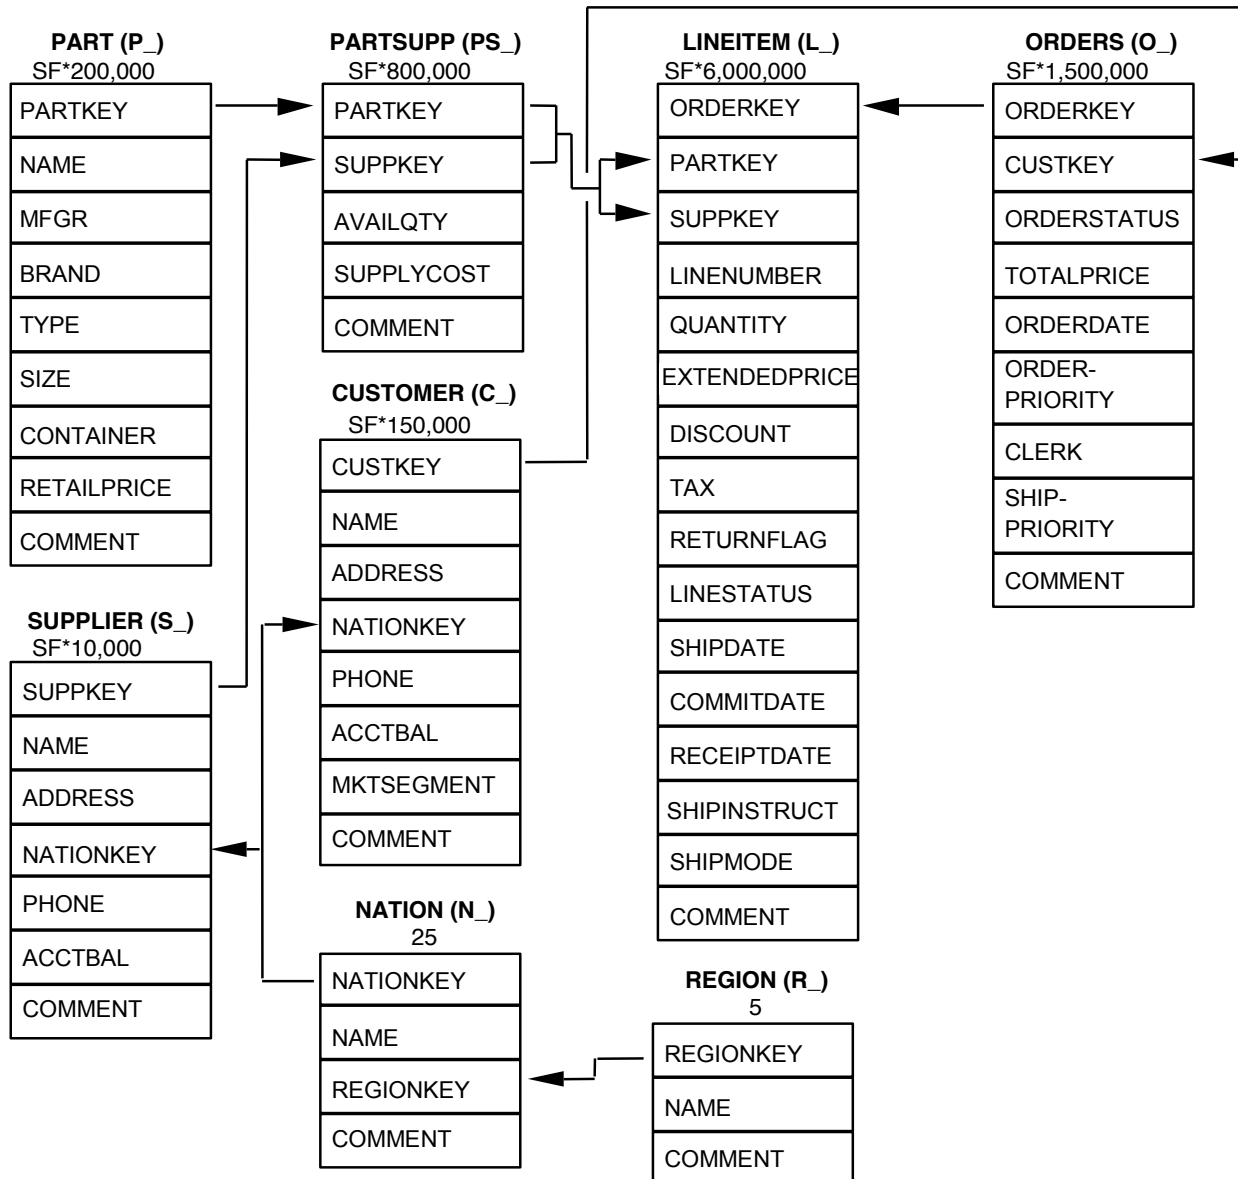


What's the assignment?

SQL-on-Hadoop



What's the assignment?



What's the assignment?

```
select
  l_returnflag,
  l_linestatus,
  sum(l_quantity) as sum_qty,
  sum(l_extendedprice) as sum_base_price,
  sum(l_extendedprice*(1-l_discount)) as sum_disc_price,
  sum(l_extendedprice*(1-l_discount)*(1+l_tax)) as sum_charge,
  avg(l_quantity) as avg_qty,
  avg(l_extendedprice) as avg_price,
  avg(l_discount) as avg_disc,
  count(*) as count_order
from lineitem
where
  l_shipdate = 'YYYY-MM-DD'
group by l_returnflag, l_linestatus;
```

SQL query  Raw Spark program

Your task...





Questions?